

# CS60021: Scalable Data Mining

## Large Scale Machine Learning

Sourangshu Bhattacharya



Convex opti.


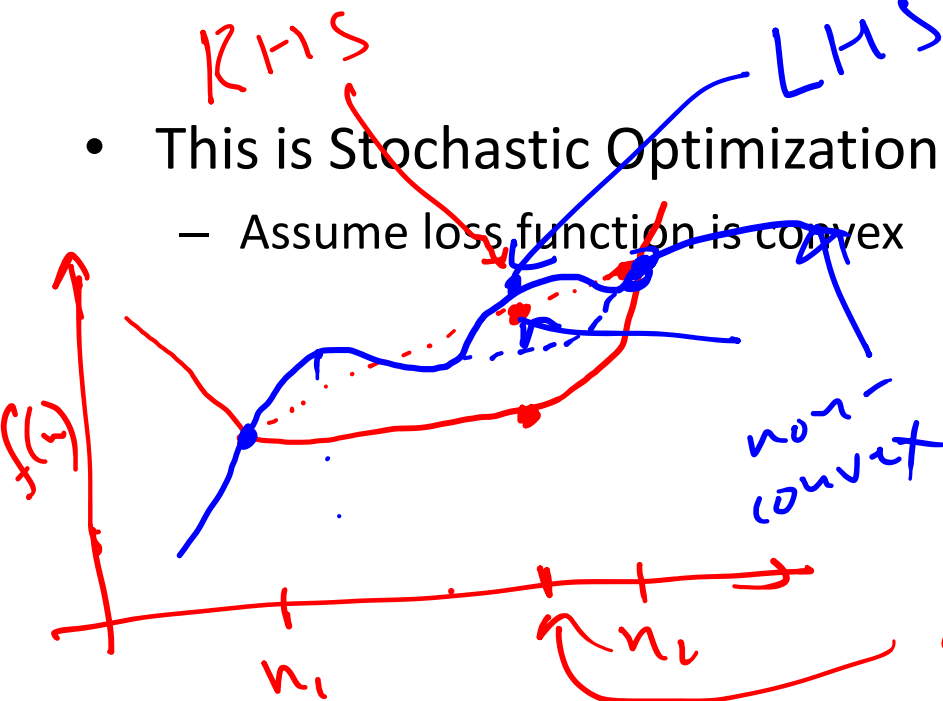
## Stochastic optimization

- Goal of machine learning :
  - Minimize expected loss

$$\min_h L(h) = \mathbf{E} [\text{loss}(h(x), y)]$$

given samples  $(x_i, y_i) \ i = 1, 2 \dots m$

- This is Stochastic Optimization
  - Assume loss function is convex



$\lambda \in [0, 1]$

$$f(\lambda n_1 + (1-\lambda)n_2)$$
$$\leq \lambda f(n_1) + (1-\lambda)f(n_2)$$

$\forall n_1, n_2 \quad f \text{ as convex}$

$$\lambda(n_1) + (1-\lambda)n_2$$

$5000$   
 $25 + 10 \leftarrow \text{grad. } w - 500$   
 $5 \times 10^6$

# Batch (sub)gradient descent for ML

10,000

- Process all examples together in each step

$$w^{(k+1)} \leftarrow w^{(k)} - \eta_t \left( \frac{1}{n} \sum_{i=1}^n \frac{\partial L(w, x_i, y_i)}{\partial w} \right)$$

where  $L$  is the regularized loss function

- ✓ Entire training set examined at each step
- Very slow when  $n$  is very large

step towards -ve grad. direction.



SGD  $\leftarrow$  stochastic grad.  
robust to training data "stochastic gradient"

## Stochastic (sub)gradient descent

- "Optimize" one example at a time
- Choose examples randomly (or reorder and choose in order)

– Learning representative of example distribution  
for  $i = 1$  to  $n$ :

$$w^{(k+1)} \leftarrow w^{(k)} - \eta_t \frac{\partial L(w, x_i, y_i)}{\partial w}$$

where  $L$  is the regularized loss function

stochastic  
optimization

$\frac{\partial L(w, x_i, y_i)}{\partial w}$

$n \times \# \text{ epochs}$

loss =

$$E_{(n,i) \sim P} [L(w, x_i, y_i)]$$
$$\nabla \text{loss} = E_{(n,i) \sim P} \left[ \frac{\partial L(w, x_i, y_i)}{\partial w} \right]$$

# Stochastic (sub)gradient descent


for  $i = 1$  to  $n$ :

$$w^{(k+1)} \leftarrow w^{(k)} - \eta_t \frac{\partial L(w, x_i, y_i)}{\partial w}$$

where  $L$  is the regularized loss function

- Equivalent to online learning (the weight vector  $w$  changes with every example)
- Convergence guaranteed for convex functions (to local minimum)

# SGD convergence



# Stochastic gradient descent

- Given dataset  $D = \{(x_1, y_1), \dots, (x_m, y_m)\}$   $\leftarrow$
- Loss function:  $L(\theta, D) = \frac{1}{N} \sum_{i=1}^N l(\theta; x_i, y_i)$   $\leftarrow$
- For linear models:  $l(\theta; x_i, y_i) = l(y_i, \theta^T \phi(x_i))$   $\leftarrow$
- Assumption  $\underline{D}$  is drawn IID from some distribution  $\mathcal{P}$ .
- Problem:

$$\min_{\theta} L(\theta, D) \quad \leftarrow$$



$$u_t \rightarrow \underline{\underline{E(l(\bar{\theta}))}} \rightarrow E(l(\theta^*))$$

## Stochastic gradient descent

- Input:  $D$
- Output:  $\bar{\theta}$

$$\frac{1}{N}$$

$$(u_t, \gamma_t) \sim \text{I.I.D.} (D)$$

### Algorithm:

- Initialize  $\theta^0$
- For  $t = 1, \dots, T$ 

$$\theta^{t+1} = \theta^t - \eta_t \nabla_{\theta} l(y_t, \theta^T \phi(x_t))$$

$$\bar{\theta} = \frac{\sum_{t=1}^T \eta_t \theta^t}{\sum_{t=1}^T \eta_t}$$

$$\bar{\theta} = \sum_{t=1}^T w_t \theta^t \quad w_t = \frac{\eta_t}{\sum \eta_t}$$

# SGD convergence

- Expected loss:  $s(\theta) = E_{\mathcal{P}}[l(y, \theta^T \phi(x))]$
- Optimal Expected loss:  $\bar{s}^* = s(\theta^*) = \min_{\theta} s(\theta)$
- Convergence:

$$E_{\bar{\theta}}[s(\bar{\theta})] - s^* \leq \frac{R^2 + L^2 \sum_{t=1}^T \eta_t^2}{2 \sum_{t=1}^T \eta_t}$$

- Where:  $R = \|\theta^0 - \theta^*\|$
- $L = \max \|\nabla l(y, \theta^T \phi(x))\|$

$$\sum \eta_t^2 < \infty$$

$$\sum \eta_t \rightarrow \infty$$

$$s(\bar{\theta})$$

$$\sum \eta_t^2$$

$$\sum \eta_t$$

$$\frac{1}{n} (l(n_1) + l(n_2) + \dots + l(n_n)) \rightarrow \sum \frac{1}{t^2}$$

$$\frac{1}{t}$$

$$T$$

# SGD convergence proof

- Define  $r_t = \|\theta^t - \theta^*\|$  and  $g_t = \nabla_{\theta} l(y_t, \theta^T \phi(x_t))$
- $r_{t+1}^2 = r_t^2 + \eta_t^2 \|g_t\|^2 - 2\eta_t (\theta^t - \theta^*)^T g_t$
- Taking expectation w.r.t  $\mathcal{P}, \bar{\theta}$  and using  $s^* - s(\theta^t) \geq g_t^T (\theta^* - \theta^t)$ , we get:

$$E_{\bar{\theta}}[r_{t+1}^2 - r_t^2] \leq \eta_t^2 L^2 + 2\eta_t (s^* - E_{\bar{\theta}}[s(\theta^t)])$$

- Taking sum over  $t = 1, \dots, T$  and using

$$E_{\bar{\theta}}[r_{t+1}^2 - r_0^2] \leq L^2 \sum_{t=0}^{T-1} \eta_t^2 + 2 \sum_{t=0}^{T-1} \eta_t (s^* - E_{\bar{\theta}}[s(\theta^t)])$$

# SGD convergence proof

- Using convexity of  $s$ :

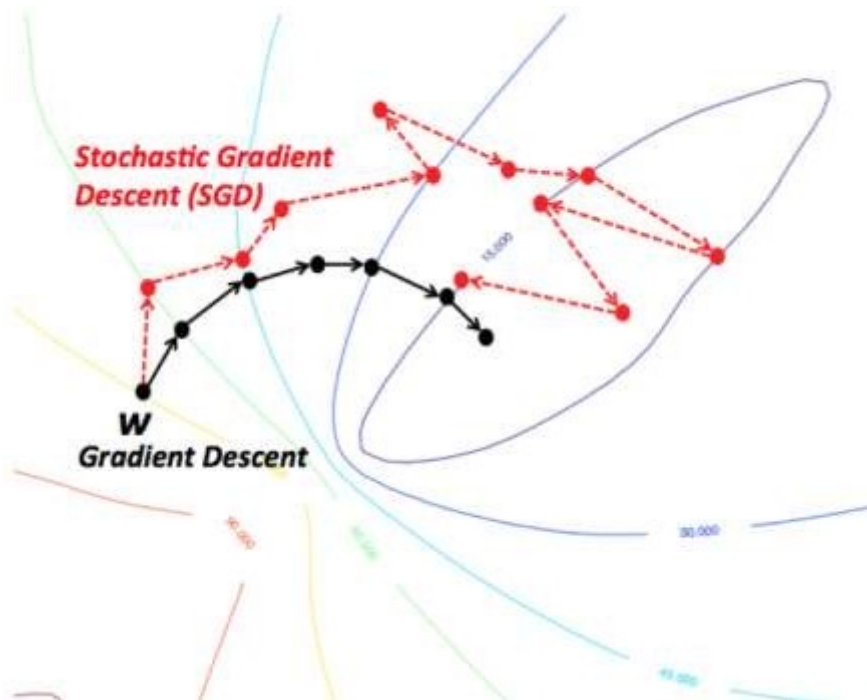
$$\left( \sum_{t=0}^{T-1} \eta_t \right) E_{\bar{\theta}} [s(\bar{\theta})] \leq E_{\bar{\theta}} \left[ \sum_{t=0}^{T-1} \eta_t s(\theta^t) \right]$$

- Substituting in the expression from previous slide:

$$E_{\bar{\theta}} [r_{t+1}^2 - r_0^2] \leq L^2 \sum_{t=0}^{T-1} \eta_t^2 + 2 \sum_{t=0}^{T-1} \eta_t (s^* - E_{\bar{\theta}} [s(\bar{\theta})])$$

- Rearranging the terms proves the result.

# The fluctuation : Batch vs SGD



<https://wikidocs.net/3413>

Batch gradient descent converges to the minimum of the basin the parameters are placed in and the **fluctuation is small**.

**SGD's fluctuation is large** but it enables to jump to new and potentially better local minima.

However, this ultimately complicates convergence to the exact minimum, as SGD will keep overshooting

## SGD - Issues

- Convergence very sensitive to learning rate ( $\eta_t$ ) (oscillations near solution due to probabilistic nature of sampling)
  - Might need to decrease with time to ensure the algorithm converges eventually
- Basically – SGD good for machine learning with large data sets!

## Mini-batch SGD

- Stochastic – 1 example per iteration
- Batch – All the examples!
- Mini-batch SGD:
  - Sample  $m$  examples at each step and perform SGD on them
- Allows for parallelization, but choice of  $m$  based on heuristics

# Example: Text categorization

- **Example by Leon Bottou:**
  - **Reuters RCV1** document corpus
    - Predict a category of a document
      - One **vs.** the rest classification
  - **$n = 781,000$**  training examples (documents)
  - 23,000 test examples
  - **$d = 50,000$**  features
    - One feature per word
    - Remove stop-words
    - Remove low frequency words



# Example: Text categorization


- Questions:

- (1) Is **SGD** successful at minimizing  $f(\mathbf{w}, \mathbf{b})$ ?
- (2) How quickly does **SGD** find the min of  $f(\mathbf{w}, \mathbf{b})$ ?
- (3) What is the error on a test set?

	<i>Training time</i>	<i>Value of <math>f(\mathbf{w}, \mathbf{b})</math></i>	<i>Test error</i>
Standard SVM	23,642 secs	0.2275	6.02%
“Fast SVM”	66 secs	0.2278	6.03%
<b>SGD SVM</b>	1.4 secs	0.2275	6.02%

- (1) SGD-SVM is successful at minimizing the value of  $f(\mathbf{w}, \mathbf{b})$
- (2) SGD-SVM is super fast
- (3) SGD-SVM test set error is comparable

# Optimization “Accuracy”




Optimization quality:  $|f(w, b) - f(w^{opt}, b^{opt})|$

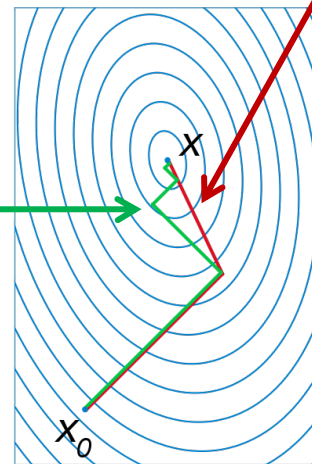
For optimizing  $f(w, b)$  within reasonable quality SGD-SVM is super fast

# SGD vs. Batch Conjugate Gradient

**SGD** on full dataset vs. **Conjugate Gradient** on a sample of  $n$  training examples



Theory says: Gradient descent converges in linear time  $k$ . Conjugate gradient converges in  $\sqrt{k}$ .  $k$ ... condition number



**Bottom line:** Doing a simple (but fast) SGD update many times is better than doing a complicated (but slow) CG update a few times

# Practical Considerations


- Need to choose learning rate  $\eta$  and  $t_0$

$$w_{t+1} \leftarrow w_t - \frac{\eta_0}{t + t_0} \left( w_t + C \frac{\partial L(x_i, y_i)}{\partial w} \right)$$

- Leon suggests:

- Choose  $t_0$  so that the expected initial updates are comparable with the expected size of the weights
- Choose  $\eta$ :
  - Select a **small subsample**
  - Try various rates  $\eta$  (e.g., 10, 1, 0.1, 0.01, ...)
  - Pick the one that most reduces the cost
  - Use  $\eta$  for next 100k iterations on the full dataset

# Learning rate comparison



# Practical Considerations

- **Sparse Linear SVM:**

- **Feature vector  $\mathbf{x}_i$  is sparse (contains many zeros)**

- Do not do:  $\mathbf{x}_i = [0, 0, 0, 1, 0, 0, 0, 0, 5, 0, 0, 0, 0, 0, \dots]$
- But represent  $\mathbf{x}_i$  as a sparse vector  $\mathbf{x}_i = [(4, 1), (9, 5), \dots]$

- **Can we do the SGD update more efficiently?**

$$w \leftarrow w - \eta \left( w + C \frac{\partial L(x_i, y_i)}{\partial w} \right)$$

- **Approximated in 2 steps:**

$$w \leftarrow w - \eta C \frac{\partial L(x_i, y_i)}{\partial w} \quad \text{cheap: } \mathbf{x}_i \text{ is sparse and so few coordinates } \mathbf{j} \text{ of } \mathbf{w} \text{ will be updated}$$

$$w \leftarrow w(1 - \eta) \quad \text{expensive: } \mathbf{w} \text{ is not sparse, all coordinates need to be updated}$$

# Practical Considerations

## ■ **Solution 1:** $\mathbf{w} = \mathbf{s} \cdot \mathbf{v}$

- Represent vector  $\mathbf{w}$  as the product of scalar  $\mathbf{s}$  and vector  $\mathbf{v}$
- Then the update procedure is:
  - $(1) \mathbf{v} = \mathbf{v} - \eta C \frac{\partial L(x_i, y_i)}{\partial \mathbf{w}}$
  - $(2) \mathbf{s} = \mathbf{s}(1 - \eta)$

## • **Solution 2:**

- Perform only step **(1)** for each training example
- Perform step **(2)** with lower frequency and higher  $\eta$

**Two step update procedure:**

$$(1) w \leftarrow w - \eta C \frac{\partial L(x_i, y_i)}{\partial w}$$

$$(2) w \leftarrow w(1 - \eta)$$

# Practical Considerations

- **Stopping criteria:**

## How many iterations of SGD?

- **Early stopping with cross validation**

- Create a validation set
- Monitor cost function on the validation set
- Stop when loss stops decreasing

- **Early stopping**

- Extract two disjoint subsamples **A** and **B** of training data
- Train on **A**, stop by validating on **B**
- Number of epochs is an estimate of  $k$
- Train for  $k$  epochs on the full dataset



# **ACCELERATED GRADIENT DESCENT**

# Stochastic gradient descent

Idea: Perform a parameter update for each training example  $x(i)$  and label  $y(i)$

Update:  $\theta = \theta - \eta \cdot \nabla_{\theta} J(\theta; x(i), y(i))$

Performs redundant computations for large datasets

# Momentum gradient descent

- Idea: Overcome ravine oscillations by momentum

- 

Update:


- $V_t = \gamma V_{t-1} + \eta \cdot \nabla_{\theta} J(\theta)$

- $\theta = \theta - V_t$

SGD




SGD with  
momentum



# Why Momentum Really Works

The momentum term **reduces updates for dimensions whose gradients change directions.**




The momentum term **increases for dimensions whose gradients point in the same directions.**

Demo : <http://distill.pub/2017/momentum/>

# Nesterov accelerated gradient

- However, a ball that rolls down a hill, blindly following the slope, is highly unsatisfactory.
- We would like to have a smarter ball that has a notion of where it is going so that it knows to slow down before the hill slopes up again.
- **Nesterov accelerated gradient** gives us a way of it.

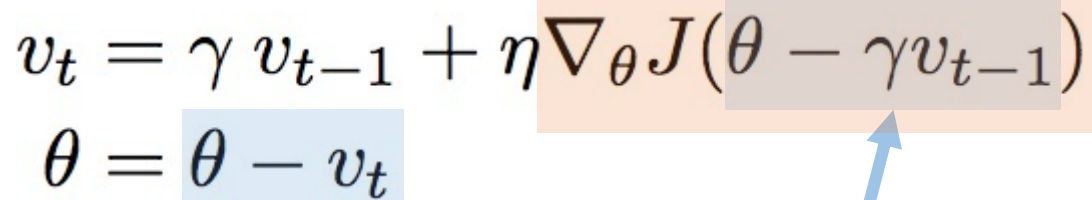
# Nesterov accelerated gradient

$$v_t = \gamma v_{t-1} + \eta \nabla_{\theta} J(\theta - \gamma v_{t-1})$$
$$\theta = \theta - v_t$$


Approximation of the next position of the parameters(predict)


# Nesterov accelerated gradient

Approximation of the next position of the parameters' gradient(**correction**)

$$v_t = \gamma v_{t-1} + \eta \nabla_{\theta} J(\theta - \gamma v_{t-1})$$
$$\theta = \theta - v_t$$


Approximation of the next position of the parameters(**predict**)

# Nesterov accelerated gradient



Blue line : predict

Red line : correction

Green line : accumulated gradient


Approximation of the next position  
of the parameters'  
gradient(**correction**)

$$v_t = \gamma v_{t-1} + \eta \nabla_{\theta} J(\theta - \gamma v_{t-1})$$
$$\theta = \theta - v_t$$

Approximation of the next position  
of the parameters(**predict**)



# Nesterov accelerated gradient



Blue line : predict

Red line : correction


Green line : accumulated gradient

Approximation of the next position  
of the parameters'  
gradient(**correction**)

$$v_t = \gamma v_{t-1} + \eta \nabla_{\theta} J(\theta - \gamma v_{t-1})$$
$$\theta = \theta - v_t$$

Approximation of the next position  
of the parameters(**predict**)

# Nesterov accelerated gradient



Blue line : predict

Red line : correction

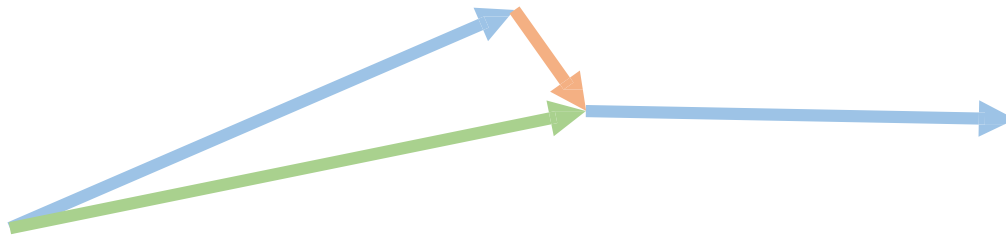
Green line : accumulated gradient

Approximation of the next position  
of the parameters'  
gradient(**correction**)

$$v_t = \gamma v_{t-1} + \eta \nabla_{\theta} J(\theta - \gamma v_{t-1})$$
$$\theta = \theta - v_t$$

Approximation of the next position  
of the parameters(**predict**)

# Nesterov accelerated gradient



Blue line : predict

Red line : correction


Green line : accumulated gradient

Approximation of the next position of  
the parameters' gradient (correction)

$$v_t = \gamma v_{t-1} + \eta \nabla_{\theta} J(\theta - \gamma v_{t-1})$$
$$\theta = \theta - v_t$$

Approximation of the next position  
of the parameters (predict)

# Nesterov accelerated gradient



Blue line : predict

Red line : correction

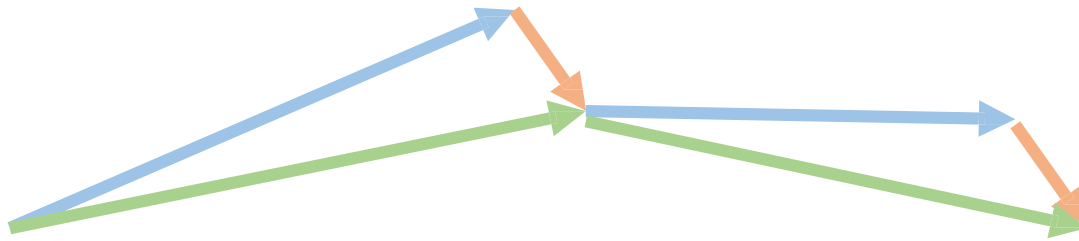
Green line : accumulated gradient

Approximation of the next position  
of the parameters'  
gradient(**correction**)

$$v_t = \gamma v_{t-1} + \eta \nabla_{\theta} J(\theta - \gamma v_{t-1})$$
$$\theta = \theta - v_t$$

Approximation of the next position  
of the parameters(**predict**)

# Nesterov accelerated gradient



Blue line : predict

Red line : correction

Green line : accumulated gradient

Approximation of the next position  
of the parameters'  
gradient(**correction**)

$$v_t = \gamma v_{t-1} + \eta \nabla_{\theta} J(\theta - \gamma v_{t-1})$$
$$\theta = \theta - v_t$$

Approximation of the next position  
of the parameters(**predict**)

# Nesterov accelerated gradient

- This anticipatory update **prevents** us from **going too fast** and **results in increased responsiveness**.
- Now , we can adapt our updates to the slope of our error function and **speed up SGD** in turn.

# AdaGrad

Adapts the learning rate to the parameters

- - Smaller updates (i.e. low learning rates) for parameters associated with frequently occurring features

larger updates (i.e. high learning rates) for parameters associated with infrequent features

Update:

$$\theta_{t+1,i} = \theta_{t,i} - \frac{\eta}{\sqrt{G_{t,ii} + \epsilon}} \cdot g_{t,i}.$$

# RMSprop


- Idea: Use the second moment of gradient vector to estimate the magnitude of update in a given direction.
- Update:
  - $E[g^2]_t = 0.9 E[g^2]_{t-1} + 0.1 g_t^2$
  - $\Delta\theta_t = - \eta / \sqrt{(E[g^2]_t + \epsilon)} \odot g_t$




# ADAM (Adaptive moment)

- Idea: In addition to storing an exponentially decaying average of past squared gradients like RMSprop, Adam also keeps an exponentially decaying average of past gradients.
- Updates:
  - $m_t = \beta_1 m_{t-1} + (1 - \beta_1) g_t$
  - $v_t = \beta_2 v_{t-1} + (1 - \beta_2) g_t^2$
  - $\hat{m}_t = m_t / (1 - \beta_1^t)$
  - $\hat{v}_t = v_t / (1 - \beta_2^t)$
  - $\vartheta_{t+1} = \vartheta_t - ( \eta / ( \sqrt{\hat{v}_t} + \epsilon ) ) \hat{m}_t$


# Visualization



# Visualization



# Enhancements comparison



# Summary

- There are two main ideas at play:
  - **Momentum** : Provide consistency in update directions by incorporating past update directions.
  - **Adaptive gradient** : Scale the scale updates to individual variables using the second moment in that direction.
  - This also relates to adaptively altering step length for each direction.

# References:

- SGD proof by Yuri Nesterov.
- MMDS <http://www.mmds.org/>
- *Blog of Sebastian Ruder* <http://ruder.io/optimizing-gradient-descent/>
- *Learning rate comparison* <https://towardsdatascience.com/learning-rate-schedules-and-adaptive-learning-rate-methods-for-deep-learning-2c8f433990d1>